

## Assignment -2

### Data Visualization and Pre-processing

Assignment Date	21 September 2022
Student Name	P. Bharathkumar
Student Roll Number	510119104003
Maximum Marks	2 Marks

#### Question-1:

##### Download the dataset:

The dataset "Churn\_Modelling.csv" was downloaded Successfully

#### Question-2:

##### Load the Dataset:

```
[ ] #load the dataset
```

```
[1] from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[84] import pandas as pd
import numpy as np
import sklearn
import matplotlib.pyplot as plt
```

```
[3] data = pd.read_csv("/content/drive/MyDrive/Churn_Modelling.csv")
```

```
[4] data.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

#### Question-3:

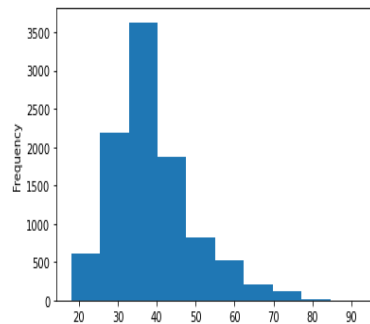
##### Perform Below Visualization:

##### Univariate Analysis

```
[5] #Univariate Analysis for Numerical data
```

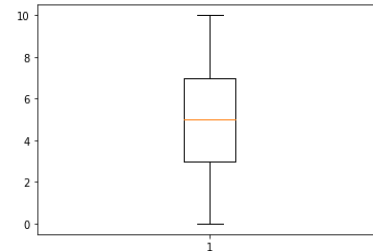
```
[6] #Histogram
data['Age'].plot(kind='hist')
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f65a0462590>



```
[9] #Box Plot
plt.boxplot(data['Tenure'])
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x7f659f8e02d0>,
<matplotlib.lines.Line2D at 0x7f659f8e0810>],
'caps': [<matplotlib.lines.Line2D at 0x7f659f8e0d50>,
<matplotlib.lines.Line2D at 0x7f659f8e62d0>],
'boxes': [<matplotlib.lines.Line2D at 0x7f659f959d10>],
'medians': [<matplotlib.lines.Line2D at 0x7f659f8e6850>],
'fliers': [<matplotlib.lines.Line2D at 0x7f659f8e6d90>],
'means': []}
```

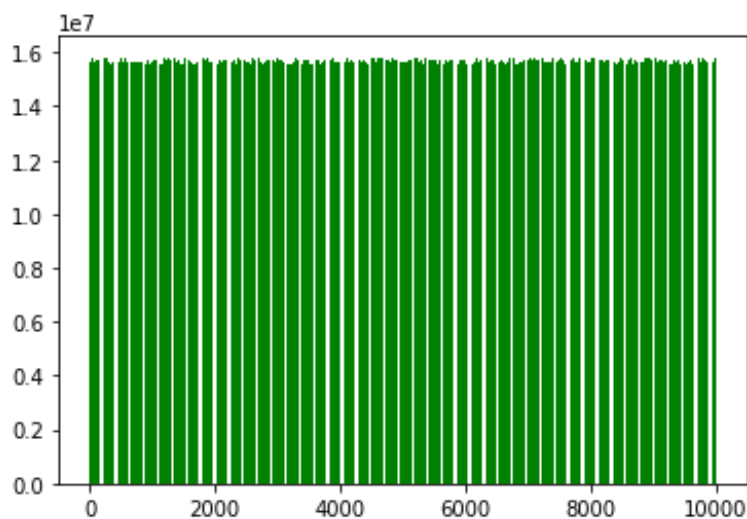


## [ ] #Univariate Analysis for Categorical Data

```
[14] #Bar Chart
df = pd.DataFrame(data)

X = list(df.iloc[:, 0])
Y = list(df.iloc[:, 1])
plt.bar(X, Y, color='g')
```

<BarContainer object of 10000 artists>




## Bivariate Analysis

✓  
19s

```
[23] #Bivariate Analysis for Categorical Data  
#Stacked Bar chart  
plt.bar(data['Gender'], data['Geography'], color='r')
```


<BarContainer object of 10000 artists>



✓  
0s

```
[21] #Bivariate Analysis for Numerical Data  
plt.scatter(data['Age'], data['CreditScore'], color='b')
```

<matplotlib.collections.PathCollection at 0x7f6589f606d0>




## Multivariate Analysis

✓  
16s

```
#Multivariate Analysis for 2 Numerical and 1 Categorical Data  
#Scatter Plot  
import seaborn as sns  
sns.catplot(data=data, x="Age", y="CreditScore", hue="Gender")
```


<seaborn.axisgrid.FacetGrid at 0x7f657aab5d90>



✓  
1s

```
[32] #Multivariate Analysis for 2 Categorical and 1 Numerical Data  
#Box Plot  
sns.catplot(data=data, x="Geography", y="Age", hue="Gender", kind="box")
```

<seaborn.axisgrid.FacetGrid at 0x7f6575c43490>



#### Question-4:

Perform Descriptive Statistics on the dataset:

```
[ ] #Perform Descriptive Statistics on the Dataset
```

```
data.mean()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:
"""Entry point for launching an IPython kernel.
RowNumber      5.000500e+03
CustomerId      1.569094e+07
CreditScore     6.505288e+02
Age             3.892180e+01
Tenure          5.012800e+00
Balance         7.648589e+04
NumOfProducts  1.530200e+00
HasCrCard       7.055000e-01
IsActiveMember  5.151000e-01
EstimatedSalary 1.000902e+05
Exited          2.037000e-01
dtype: float64
```

```
[34] data.median()

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:
"""Entry point for launching an IPython kernel.
RowNumber      5.000500e+03
CustomerId      1.569074e+07
CreditScore     6.520000e+02
Age             3.700000e+01
Tenure          5.000000e+00
Balance         9.719854e+04
NumOfProducts  1.000000e+00
HasCrCard       1.000000e+00
IsActiveMember  1.000000e+00
EstimatedSalary 1.001939e+05
Exited          0.000000e+00
dtype: float64
```

```
[36] data.describe()


```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.00000	1.000000e+04	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000	10000.00000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.558570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	0.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	1.000000

```
[38] data.shape
```

```
(10000, 14)
```

### Question-5:

Handle the Missing values:

✓ [39] #Handling the missing values

0s

```
data.isnull().sum()
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

### Question-6:


Find the outliers and replace the outliers:

```
[ ] #Find the Outliers and replace the outliers
```

✓ [40] sns.boxplot(data['Age'])

s

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x.
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f6575aed650>
```



```
✓ [41] qnt=data.quantile(q=[0.25,0.75])  
Ds qnt
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0.25	2500.75	15628528.25	584.0	32.0	3.0	0.00	1.0	0.0	0.0	51002.1100	0.0
0.75	7500.25	15753233.75	718.0	44.0	7.0	127644.24	2.0	1.0	1.0	149388.2475	0.0

```
✓ [42] IQR = qnt.loc[0.75] - qnt.loc[0.25]  
Ds IQR
```

```
RowNumber      4999.5000  
CustomerId      124705.5000  
CreditScore      134.0000  
Age              12.0000  
Tenure           4.0000  
Balance          127644.2400  
NumOfProducts     1.0000  
HasCrCard         1.0000  
IsActiveMember    1.0000  
EstimatedSalary   98386.1375  
Exited           0.0000  
dtype: float64
```

```
✓ [43] upper_extreme = qnt.loc[0.75]+1.5*IQR  
Ds upper_extreme
```

```
RowNumber      1.499950e+04  
CustomerId      1.594029e+07  
CreditScore      9.190000e+02  
Age              6.200000e+01  
Tenure           1.300000e+01  
Balance          3.191106e+05  
NumOfProducts     3.500000e+00  
HasCrCard         2.500000e+00  
IsActiveMember    2.500000e+00  
EstimatedSalary   2.969675e+05  
Exited           0.000000e+00  
dtype: float64
```

```
✓ [44] lower_extreme = qnt.loc[0.25]-1.5*IQR  
Ds lower_extreme
```

```
RowNumber      -4.998500e+03  
CustomerId      1.544147e+07  
CreditScore      3.830000e+02  
Age              1.400000e+01  
Tenure          -3.000000e+00  
Balance          -1.914664e+05  
NumOfProducts     -5.000000e-01  
HasCrCard         -1.500000e+00  
IsActiveMember    -1.500000e+00  
EstimatedSalary   -9.657710e+04  
Exited           0.000000e+00  
dtype: float64
```

```

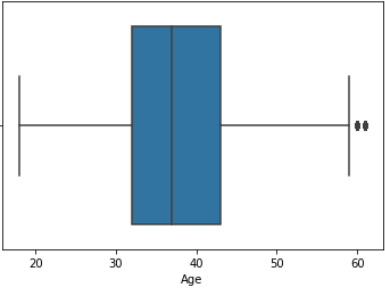
✓ [51] df2 = data[(data['Age'] < upper_extreme['Age']) & (data['Age'] > lower_extreme['Age'])]
0s

✓ [50] data.shape
1s
(10000, 14)

✓ [49] df2.shape
0s
(9589, 14)

✓ [52] sns.boxplot(df2['Age'])
0s
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x.
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f6573caad10>

```



### Question-7:

Check for Categorical columns and perform Encoding:

```

✓ [53] #Check for Categorical columns and perform encoding
0s
#Categorical are Geography and Gender
from sklearn.preprocessing import LabelEncoder

✓ [75] le=LabelEncoder()
0s
df2['Geography'] = le.fit_transform(df2['Geography'])
df2['Gender'] = le.fit_transform(df2['Gender'])

```

```
[76] df2.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	0	0	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	0	0	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	0	0	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	2	0	43	2	125510.82	1	1	1	79084.10	0



### Question-8:

Split the data into dependent and independent variables:

```
[77] #Split the data into dependent and independent variables.  
y=df2['EstimatedSalary']  
x=df2.drop(columns=['EstimatedSalary'],axis=1)
```

✓ [78] x.head()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	Exited
0	1	15634602	Hargrave	619	0	0	42	2	0.00	1	1	1	1
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	1	0
2	3	15619304	Onio	502	0	0	42	8	159660.80	3	1	0	1
3	4	15701354	Boni	699	0	0	39	1	0.00	2	0	0	0
4	5	15737888	Mitchell	850	2	0	43	2	125510.82	1	1	1	0

### Question-9:

Scale the independent variables:

```
✓ [80] #Scale the independent variables  
0s names=z.columns  
names
```

```
Index(['CreditScore', 'Geography', 'Gender', 'Age'], dtype='object')
```

```
✓ [81] from sklearn.preprocessing import scale  
0s
```

```
✓ [82] z=scale(z)  
0s z
```

```
array([[ -0.32370448, -0.90175758, -1.09674455,  0.50205394],  
       [ -0.43751069,  1.51663241, -1.09674455,  0.38636147],  
       [ -1.53418871, -0.90175758, -1.09674455,  0.50205394],  
       ...,  
       [  0.60743723, -0.90175758, -1.09674455, -0.19210091],  
       [  1.25923643,  0.30743742,  0.91178935,  0.50205394],  
       [  1.46615681, -0.90175758, -1.09674455, -1.1176407 ]])
```

```
✓ [87] z = pd.DataFrame(z, columns =names)  
0s
```



```
z.head()
```



	CreditScore	Geography	Gender	Age
0	-0.323704	-0.901758	-1.096745	0.502054
1	-0.437511	1.516632	-1.096745	0.386361
2	-1.534189	-0.901758	-1.096745	0.502054
3	0.503977	-0.901758	-1.096745	0.154977
4	2.066226	1.516632	-1.096745	0.617746



### Question-10:

Split the data into training and testing:

```
[ ] #Split The data into Training and Testing
```



0s

```
[89] from sklearn.model_selection import train_test_split
```



0s

```
[90] x_train,x_test,y_train,y_test = train_test_split(z,y,test_size=0.2)
```





0s


```
[92] x_train
```

	CreditScore	Geography	Gender	Age
1060	-0.634085	-0.901758	0.911789	-1.117641
6074	-0.437511	1.516632	-1.096745	-0.307793
7328	-1.347960	0.307437	-1.096745	-1.696103
5165	0.907472	1.516632	0.911789	0.733439
2750	-0.489241	-0.901758	-1.096745	1.196209
...	...	...	...	...
4092	1.010932	-0.901758	0.911789	0.733439
996	-0.261628	-0.901758	-1.096745	-1.580411
8690	1.186814	0.307437	-1.096745	0.964824
8099	0.131520	1.516632	-1.096745	-0.539178
1888	0.048752	1.516632	-1.096745	-0.192101




✓  y\_train

 1104 151645.96  
6334 143463.28  
7638 37577.66  
5392 43018.82  
2851 100478.60  
...  
4269 2048.55  
1037 180969.55  
9056 166896.01  
8440 36864.05  
1960 86013.96  
Name: EstimatedSalary, Length: 7671, dtype: float64

✓  [94] x\_test

	CreditScore	Geography	Gender	Age
<b>962</b>	0.772974	0.307437	0.911789	0.154977
<b>5257</b>	1.248890	1.516632	-1.096745	0.386361
<b>7515</b>	-0.841005	0.307437	-1.096745	-0.654871
<b>6844</b>	0.959202	-0.901758	-1.096745	-0.886256
<b>4102</b>	-0.996196	1.516632	-1.096745	0.386361
...	...	...	...	...
<b>60</b>	0.379825	0.307437	-1.096745	-1.233333
<b>5555</b>	0.503977	-0.901758	0.911789	-0.076408
<b>5112</b>	1.704115	1.516632	-1.096745	2.237441
<b>138</b>	0.131520	-0.901758	0.911789	-0.423486
<b>4973</b>	0.328095	-0.901758	-1.096745	2.353134

1918 rows x 4 columns

✓  [95] y\_test

1002 184023.54  
5486 92914.67  
7838 132038.65  
7133 138780.89  
4281 36242.19  
...  
61 126494.82  
5797 83263.04  
5337 38941.44  
141 180427.24  
5191 706.50  
Name: EstimatedSalary, Length: 1918, dtype: float64